

SPATIAL MODEL CHECKING WITH MICROSERVICES FOR THE INTERNET-OF-THINGS

March 15, 2018

Politecnico di Milano, Italy

- Wide deployments of internet-enabled things within smart applications (i.e. smart cities)
- Complex relations between locations of things with respect to their environment may arise
- Things change location; location may affect satisfaction of system requirements

- Wide deployments of internet-enabled things within smart applications (i.e. smart cities)
- Complex relations between locations of things with respect to their environment may arise
- Things change location; location may affect satisfaction of system requirements

How can we **enable runtime verification** of requirements of space-dependent systems of internet-of-things?

- Wide deployments of internet-enabled things within smart applications (i.e. smart cities)
- Complex relations between locations of things with respect to their environment may arise
- Things change location; location may affect satisfaction of system requirements

How can we **enable runtime verification** of requirements of space-dependent systems of internet-of-things?

- Properties expressing arbitrary relations in space are known upfront
- We seek formal assurances that properties hold in a given space-things configuration

Verify spatial properties on configurations of city space populated with locations of things at given point

Key intuition: Space is composed of relations between entities

- Evaluation model is a discrete graph structure representing a topological view of space
- Serves as the critical abstraction step
- Enables formal specification and verification techniques

EXPRESSING & VERIFYING PROPERTIES

A **spatial logic** can express complex properties of space and enable automated verification. For example:

Example

From the location where a taxi is near a gas station, is there a way to go to a metro station going through –an arbitrary number of– bus stops?

Example

From which locations where taxis are near a pharmacy, we can reach a bar through metro stations and bus stops, without going near a hospital?

We can support logic-based reasoning

- Formal specification of properties
- Spatial properties may be non trivial

SPATIAL EVALUATION MODEL

1. Generate a **closure space** from the accessibility graph
2. Equip with a **valuation function** associating points with atomic propositions (POIs, taxis etc)
3. Big-step semantics: **truth values of spatial relations** or **points where spatial relations hold**

Spatial Logic for Closure Spaces¹

$$\tau ::= p \mid \top \mid \neg\tau \mid \tau \wedge \tau \mid \mathcal{C} \tau \mid \tau S \tau$$

- Closure **"one step"** modality
- **Surrounds** modality

¹Vincenzo Ciancia et al. "Specifying and verifying properties of space". In: *Theoretical Computer Science*. Springer, 2014, pp. 222–235.

MODELLING COMPLEX PROPERTIES IN SPACE²: REACHABILITY

Notion of path in a closure space

$$\phi \mathcal{R} \psi \stackrel{\text{def}}{=} \neg((\neg\psi) \mathcal{S} (\neg\phi))$$

Constrain the traversed path

$$\phi \mathcal{T} \psi \stackrel{\text{def}}{=} \phi \wedge ((\phi \vee \psi) \mathcal{R} \psi)$$

Predicate on entities encountered on the traversed path

$$\phi \mathcal{R}(\psi) \zeta \stackrel{\text{def}}{=} \phi \mathcal{T} ((\psi \mathcal{T} \zeta) \wedge (\psi \mathcal{T} \phi))$$

Nearness as nested applications of closure operator

$$\mathcal{N}_n \phi \stackrel{\text{def}}{=} C_n C_{n-1} \dots \phi$$

²Christos Tsigkanos, Timo Kehrer, and Carlo Ghezzi. “Modeling and Verification of Evolving Cyber-Physical Spaces”. In: *ACM SIGSOFT Symposium on the Foundations of Software Engineering*. 2017, pp. 38–48.

Example

From the location where a taxi is near a gas station, is there a way to go to a metro station going through –an arbitrary number of– bus stops?

$\text{taxi } \mathcal{R}(\text{transportbusstop}) \text{ transportsubway} \wedge \mathcal{N}_n \text{transportfuel}.$

Example

From which locations where taxis are near a pharmacy, we can reach a bar through metro stations and bus stops, without going near a hospital?

$\text{taxi } \mathcal{R}(\text{transportbusstop} \vee \text{transportsubway} \wedge !(\mathcal{N}_n \text{healthhospital}))$
 $\text{foodbar} \wedge \mathcal{N}_n \text{healthpharmacy}.$

Intuition:

- Trajectory datasets can provide movement data points
- Points-of-interest datasets can provide static graphs

Combining both

Movement data points over time, over a closure space

T-Drive trajectory dataset³

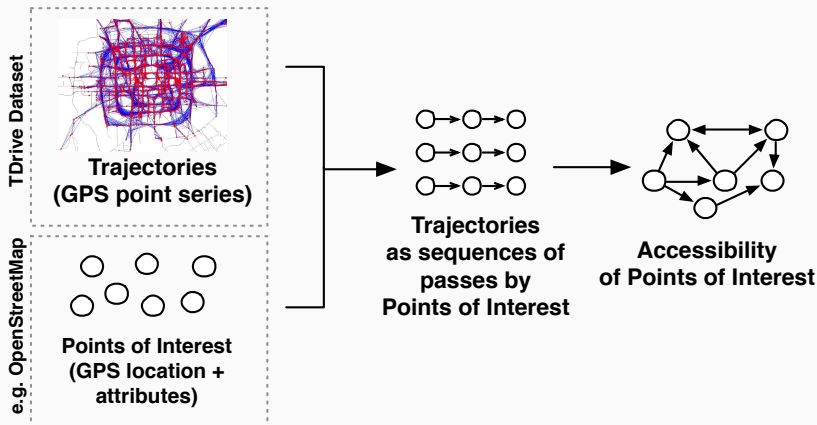
- trajectories of 10k taxis over one week
- 15M data points
- 9M kilometers total distance

OpenStreetMap points-of-interest in Beijing

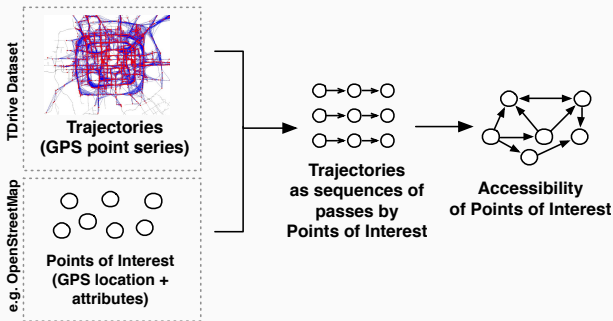
- Various landmarks, public services, shops, etc
- e.g. 11741 POIs

³Jing Yuan et al. “Driving with knowledge from the physical world”. In: *Proceedings of the 17th ACM SIGKDD international conference on Knowledge discovery and data mining*. ACM. 2011, pp. 316–324, Jing Yuan et al. “T-drive: driving directions based on taxi trajectories”. In: *Proceedings of the 18th SIGSPATIAL International conference on advances in geographic information systems*. ACM. 2010, pp. 99–108.

OBTAINING A TOPOLOGICAL MODEL: PROCESS



OBTAINING A TOPOLOGICAL MODEL: PROCESS



1. T-Drive trajectory dataset
2. Points-of-Interest
3. Discover points that trajectories pass through

Key idea: if an entity's trajectory passes from a point to another, then these points are connected

- We obtain connectivity of points as a graph structure
- Nodes are Points-of-Interest
- Edges reflect “accessibility” of a node from another

Discrete presences of taxis in Beijing over time

- Taxis move between various points of interest
- Can consider them as position *updates* to the global topological space of Beijing

Intuition

Every taxi presence triggers a check of a global spatial property

EVALUATION DATASET: A PROPERTY


“Always from where taxis are near metro stations, there is a way to go to gas stations going through –an arbitrary number of– bus stops or bars ”

Formulation

$$\begin{aligned} & (\text{taxi} \wedge \mathcal{C}\text{transportsubway}) \wedge \neg((\neg(((\text{transportbusstop} \vee \text{foodbar}) \wedge \\ & \neg((\neg(\text{transportfuel}))\mathcal{S}(\neg((\text{transportbusstop} \vee \text{foodbar}) \vee \\ & (\text{transportfuel})))))) \wedge ((\text{transportbusstop} \vee \text{foodbar}) \wedge \neg((\neg(\text{taxi} \wedge \\ & \mathcal{C}\text{transportsubway}))\mathcal{S}(\neg((\text{transportbusstop} \vee \text{foodbar}) \vee (\text{taxi} \wedge \\ & \mathcal{C}\text{transportsubway}))))))\mathcal{S}(\neg((\text{taxi} \wedge \mathcal{C}\text{transportsubway}) \vee \\ & (((\text{transportbusstop} \vee \text{foodbar}) \wedge \\ & \neg((\neg(\text{transportfuel}))\mathcal{S}(\neg((\text{transportbusstop} \vee \text{foodbar}) \vee \\ & (\text{transportfuel})))))) \wedge ((\text{transportbusstop} \vee \text{foodbar}) \wedge \neg((\neg(\text{taxi} \wedge \\ & \mathcal{C}\text{transportsubway}))\mathcal{S}(\neg((\text{transportbusstop} \vee \text{foodbar}) \vee (\text{taxi} \wedge \\ & \mathcal{C}\text{transportsubway})))))))))) \end{aligned}$$

(i.e. not trivial -;)

BEIJING COMPOSITE DATASET: LUNCHTIME REQUESTS WORKLOAD



We deployed a model checker (computation logic) and the current topology (static state) as a RESTful microservice.

Three alternative architectures:

- **Lambda Functions:** 3Gb RAM (vCPU quotas are assigned accordingly)
- **Single Instance of a Big VM:** Standard F64s_v2 (64 vcpus, 128 GB memory)
- **Autoscaling Group:** 5 to 20 t1.micro VMs (1Gb Ram, 1vCPU each)


Positions tracker (dynamic state) deployed as another standalone microservice

- **One hour window** of the Beijing dataset as workload (536 data points), the beginning of a daily peak
- **Time multipliers:** To simulate events that happen closer to each other.
 - 5x / 10x / 20x / 40x / 60x time multipliers.
- Run the experiment combining the different time multipliers with the three alternative deployments

Service Level Agreement

Defined a SLA of 30 seconds/call maximum for responding to model checking requests

EXPERIMENTAL RESULTS




- The Autoscaling group had the **worst performance**. Requests timed out (504) with 20x time or more.
- Lambda held an almost **constant performance** (12s/call) for all workloads
- The BigVM (**baseline**) performed better (9-12s) up to 20x, then started to decrease performance.

EXPERIMENTAL RESULTS


Computation Time per Call (seconds)

Median





Wait time per Call (seconds)

Median (Logarithmic scale)



EXPERIMENTAL RESULTS



EXPERIMENTAL RESULTS: DISCUSSION

- Lambda functions kept a **constant performance** even under dense workloads, being more reactive and scalable.
- The BigVM used as a baseline performed well for medium workloads, but not against peaks or more dense workloads (**e.g. workloads unknown beforehand**).
- The autoscaling group did not seem suitable, but more experimentation is needed (fine-tuning parameters, scaling containers?).
- Cost analysis to come (if the workload is more or less known, lambda could become more costly in comparison)